

ATAL Sponsored online five-Days Faculty Development Programme
on
“Data Analytics using Python/R programming”

Organized by
Department of CSE, NIT Puducherry

 **AMRITA** School of
Engineering
R packages for data science

Presented by: Siddharth R

Department of Computer Science and Engineering
Amrita School of Engineering, Amrita Vishwa Vidyapeetham,
Chennai Campus



August 27, 2021



Agenda



dplyr package

caret package

Section 1



Manipulating Data frames

- ▶ The data frame is a key data structure in statistics and in R.
- ▶ There is one observation per row and each column represents a variable, a measure, feature, or characteristic of that observation
- ▶ we have already discussed use of `[]` and `$` operators to extract subsets of data frames
- ▶ The *dplyr* package provide a highly optimized set of routines specifically for dealing with data frames.

dplyr package

- ▶ The *dplyr* package was developed by Hadley Wickham of RStudio
- ▶ The dplyr package **does not** provide any “new” functionality to R per se
- ▶ Everything *dplyr* does could already be done with base R, but it greatly simplifies existing functionality in R.
- ▶ dplyr package provides a “grammar” (in particular, verbs) for data manipulation and for operating on data frames.
- ▶ The *dplyr* functions are very fast

Basic grammar in *dplyr*

- ▶ **select** - return a subset of the columns of a data frame, using a flexible notation
- ▶ **filter** - extract a subset of rows from a data frame based on logical conditions
- ▶ **arrange** - reorder rows of a data frame
- ▶ **rename** - rename variables in a data frame
- ▶ **% > %** : the “pipe” operator is used to connect multiple verb actions together into a pipeline

Common characteristics

All of the functions that we will discuss have a few common characteristics

- ▶ The first argument is a data frame.
- ▶ The subsequent arguments describe what to do with the data frame specified in the first argument, and you can refer to columns in the data frame directly without using the \$ operator
- ▶ The return result of a function is a new data frame
- ▶ Data frames must be properly formatted and annotated - there should be one observation per row, and each column should represent a feature or characteristic of that observation.

Installing and loading dplyr package

- ▶ `install.packages("dplyr")` - To install the package
- ▶ `library(dplyr)` - load it into your R session
- ▶ You may get some warnings when the package is loaded because there are functions in the dplyr package that have the same name as functions in other packages.

select

- ▶ The `select()` function can be used to select **columns** of a data frame that you want to focus on.
- ▶ Often you'll have a large data frame containing "all" of the data, but any given analysis might only use a subset of variables or observations.
- ▶ The `select()` function allows you to get the few columns you might need.
- ▶ Note that the `:` normally cannot be used with names or strings, but inside the `select()` function you can use it to specify a range of variable names.
- ▶ You can also omit variables using the `select()` function by using the negative sign.

Additional Options in select

- ▶ `starts_with()` = select all columns that start with the character string
- ▶ `ends_with()` = Select columns that end with a character string
- ▶ `contains()` = Select columns that contain a character string
- ▶ `one_of()` = Select columns names that are from a group of names

filter

- ▶ The `filter()` function is used to extract subsets of rows from a data frame

```
> chic.f <- filter(chicago, pm25tmean2 > 30)
> head(chic.f)
  city tmpd dtpd      date pm25tmean2 pm10tmean2 o3tmean2 no2tmean2
1 chic   23  21.9 1998-01-17      38.10      32.46154  3.180556  25.30000
2 chic   28  25.8 1998-01-23      33.95      38.69231  1.750000  29.37630
3 chic   55  51.3 1998-04-30      39.40      34.00000 10.786232  25.31310
4 chic   59  53.7 1998-05-01      35.40      28.50000 14.295125  31.42905
5 chic   57  52.0 1998-05-02      33.30      35.00000 20.662879  26.79861
6 chic   57  56.0 1998-05-07      32.10      34.50000 24.270422  33.99167
> summary(chic.f$pm25tmean2)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 30.05  32.12   35.04   36.63  39.53   61.50
> |
```

arrange

- ▶ Used to reorder rows of a data frame according to one of the variables
- ▶ Reordering rows of a data frame is normally a pain to do in R.
- ▶ The `arrange()` function simplifies the process quite a bit.
- ▶ Columns can be arranged in descending order too by using the special `desc()` operator.

Output Screenshot

```
> ##### ARRANGE #####
> chic.wa <- chicago[order(chicago$date),] # without arrange function
> #chic.wa <- chicago[order(as.Date(chicago$date,format="%Y-%m-%d")),]
> head(chic.wa)
  city tmpd  dptp      date pm25tmean2 pm10tmean2 o3tmean2 no2tmean2
1 chic  31.5 31.500 1987-01-01          NA      34.00000  4.250000  19.98810
2 chic  33.0 29.875 1987-01-02          NA          NA  3.304348  23.19099
3 chic  33.0 27.375 1987-01-03          NA      34.16667  3.333333  23.81548
4 chic  29.0 28.625 1987-01-04          NA      47.00000  4.375000  30.43452
5 chic  32.0 28.875 1987-01-05          NA          NA  4.750000  30.33333
6 chic  40.0 35.125 1987-01-06          NA      48.00000  5.833333  25.77233
>
>
> chic.a <- arrange(chicago, date)
> head(chic.a)
  city tmpd  dptp      date pm25tmean2 pm10tmean2 o3tmean2 no2tmean2
1 chic  31.5 31.500 1987-01-01          NA      34.00000  4.250000  19.98810
2 chic  33.0 29.875 1987-01-02          NA          NA  3.304348  23.19099
3 chic  33.0 27.375 1987-01-03          NA      34.16667  3.333333  23.81548
4 chic  29.0 28.625 1987-01-04          NA      47.00000  4.375000  30.43452
5 chic  32.0 28.875 1987-01-05          NA          NA  4.750000  30.33333
6 chic  40.0 35.125 1987-01-06          NA      48.00000  5.833333  25.77233
> chic.a <- arrange(chicago, desc(date))
> head(chic.a)
  city tmpd  dptp      date pm25tmean2 pm10tmean2 o3tmean2 no2tmean2
1 chic   35 30.1 2005-12-31    15.00000      23.5  2.531250  13.25000
2 chic   36 31.0 2005-12-30    15.05714      19.2  3.034420  22.80556
3 chic   35 29.4 2005-12-29     7.45000      23.5  6.794837  19.97222
4 chic   37 34.5 2005-12-28    17.75000      27.5  3.260417  19.28563
5 chic   40 33.6 2005-12-27    23.56000      27.0  4.468750  23.50000
6 chic   35 29.6 2005-12-26     8.40000       8.5 14.041667  16.81944
> |
```

rename

- ▶ Renaming a variable in a data frame in R is surprisingly hard to do.
- ▶ The `rename()` function is designed to make this process easier.
- ▶ The syntax inside the `rename()` function is to have the new name on the left-hand side of the `=` sign and the old name on the right-hand side.
- ▶ Another way - using `select()` function to rename
- ▶ **Exercise** - Rename a column without using dplyr package

Output Screenshot

```
> head(chicago[, 1:5],3) # Forcing the head to print only first 3 three rows of the data with first
  city tmpd  dptp      date pm25tmean2
1 chic 31.5 31.500 1987-01-01      NA
2 chic 33.0 29.875 1987-01-02      NA
3 chic 33.0 27.375 1987-01-03      NA
>
> chic.rn <- rename(chicago, dewpoint = dptp, pm25 = pm25tmean2) # Renaming two column names in the
> head(chic.rn[, 1:5],3)
  city tmpd dewpoint      date pm25
1 chic 31.5   31.500 1987-01-01   NA
2 chic 33.0   29.875 1987-01-02   NA
3 chic 33.0   27.375 1987-01-03   NA
>
> ## Another way of renaming inside the select function() but remaining columns chopped-off
> chic.rn <- select(chicago, dewpoint = dptp, pm25 = pm25tmean2)
> head(chic.rn[, 1:2],3)
  dewpoint pm25
1   31.500   NA
2   29.875   NA
3   27.375   NA
> |
```

Pipe Operator

- ▶ The pipeline operator `%>%` is very handy for stringing together multiple dplyr functions in a sequence of operations.
- ▶ Every time we wanted to apply more than one function, the sequence gets buried in a sequence of nested function calls that is difficult to read `third(second(first(x)))`
- ▶ The `%>%` operator allows you to string operations in a left-to-right fashion
`first(x) %>% second %>% third`

Agenda



Section 2



ggplot2

- ▶ ggplot2 is an R package for producing statistical, or data, graphics.
- ▶ Every ggplot2 plot has three key components:
 1. **data**
 2. A set of **aesthetic** mappings between variables in the data and visual properties
 3. At least one layer which describes how to render each observation. Layers are usually created with a **geom** function.
- ▶ data and aesthetic mappings are supplied in `ggplot()`, then layers are added on with `+`.

ggplot2(contd...)

- ▶ **faceting** - It creates tables of graphics by splitting the data into subsets and displaying the same graph for each subset.
- ▶ For a different geom function, you'd get a different type of plot.
- ▶ If you have a scatter plot with a lot of noise, it can be hard to see the dominant pattern. In this case it's useful to add a smoothed line to the plot with *geom_smooth()*

Agenda



dplyr package

caret package

Section 3



caret

- ▶ Caret is short for **C**lassification **A**nd **RE**gression **T**raining.
- ▶ Why caret? - With R having so many implementations of ML algorithms, it can be challenging to keep track of which algorithm resides in which package.
- ▶ No matter which package the algorithm resides, caret will remember that for you.

Preprocessing in caret

- ▶ **range:** Normalize values so it ranges between 0 and 1
- ▶ **center:** Subtract Mean
- ▶ **scale:** Divide by standard deviation
- ▶ **pca:** Replace with principal components
- ▶ **ica:** Replace with independent components

References

1. Roger D peng, "R Programming for Data Science", Lean Publishing, 2015
2. Hadley Wickham and Garrett Golemund, "R for Data Science", O'Reilly
3. <https://www.machinelearningplus.com/machine-learning/caret-package/>

 **Thank You..** School of
Engineering
AMRITA
VISHVA VEDICAM
UNIVERSITY